

*****WARNING: This annotation task contains short fictional scenarios describing a range of human behaviors, some of which involve distressing or sensitive themes including violence, animal cruelty, child endangerment, substance abuse, theft, deception, and references to weapons and bodily harm. These scenarios are drawn from an established research benchmark (ETHICS; Hendrycks et al., 2021) and are presented solely for the purpose of evaluating moral judgment classification. None of the scenarios describe real events. If you are uncomfortable with exposure to such content, you are under no obligation to participate, and you may withdraw at any time by entering q during the annotation session without any consequence. By proceeding with the annotation, you acknowledge that you have read this warning and consent to participate.*****

You will be annotating ethical scenarios derived from the ETHICS dataset, along with minimal edit pairs. The ETHICS dataset (Hendrycks et al., 2021) is a commonsense morality benchmark consisting of short, everyday scenarios that are labeled as either *Acceptable* or *Unacceptable* based on whether the described behavior aligns with widely shared social norms. Scenarios range from mundane actions ("I donated money to a charity") to clear moral violations ("I stole the car"), and are designed to test whether a system can distinguish socially normative behavior from behavior that most reasonable adults would consider wrong.

The document contains two python scripts that you will need to run to initiate the annotation. Before running, change the ANNOTATOR_NAME variable in the second cell to your own name, it's at the very top of the second cell. After running the two scripts, ethical scenarios will appear on the screen in pairs of two in the base of the second cell. The first scenario will be the original scenario taken directly from the ETHICS dataset, along with a pre-made annotation on whether that scenario was acceptable or unacceptable. Below, there will be a second ethical scenario, which will be the original scenario with a minor edit to the wording that will either maintain the scenario's original label, or flip the scenario's original label. Here is an example of what may be returned by the script:

```
--- Pair 1/200 (ID: X) ---  
Original: "I walked my dog"  
Label:   Acceptable  
Edited:  "I walked my cat"
```

FLIP or PRESERVE? [f/p/q]: __

After this text shows up on your screen, you will enter into the text input whether you believe the edited scenario maintained the original scenario's label (i.e. the original is acceptable and you believe the edited version is also acceptable), or if the edited scenario flipped the original scenario's label (i.e. the original is acceptable and you believe the edited version is unacceptable). The three options you can enter into the textbox are f, which indicates that you believe the label was flipped, p, which indicates that you believe the label was preserved, or q, which allows you to quit out of the annotation and saves your progress. You will repeat this for

all 200 of the dataset's contrasts, after which your responses will be saved to a CSV file. After you run the first two cells and finish your annotation, you should run the third cell which will download the CSV to your computer and allow you to send the CSV to me via email.

Please complete your annotations independently without consulting other annotators or external resources. Your individual judgment is what we are measuring. When you are finished, send your CSV file to rohan.thomas@barstowschool.org.